

K2 Think - API Guide

Overview

This document provides a Python script for testing K2-Think via an OpenAI-compatible API. The endpoint accepts OpenAI Chat Completions-style requests (`messages`) and supports streaming via `stream: true`.

Testing via cURL

Code block

```
1  curl -X POST \  
2    'https://api.k2think.ai/v2/chat/completions' \  
3    -H 'accept: application/json' \  
4    -H 'Authorization: Bearer $K2THINK_API_KEY' \  
5    -H 'Content-Type: application/json' \  
6    -d '{  
7        "model": "MBZUAI-IFM/K2-Think",  
8        "messages": [  
9            {"role": "user", "content": "hello"}  
10       ],  
11        "stream": true  
12    }'
```

Output format

K2-Think wraps thinking traces and the final answer in `<think>...</think>` and `<answer>...</answer>` tags respectively.

If you only want the final answer: extract the content inside `<answer>...</answer>`

Supported Parameters

Core request fields

- `model` (*string, required*)

Example: `"MBZUAI-IFM/K2-Think"`

- `messages` (*array, required*)

Example: `[{ "role": "user", "content": "hello" }]`

- `stream` (*boolean, optional*)

Sampling / generation controls

- `temperature` (*number, optional*)
- `top_p` (*number, optional*)
- `max_tokens` (*integer, optional*)
- `stop` (*string or array, optional*)

Works both as a single string and an array of strings.

Rate Limit

- The API can handle 20 RPM. If exceeded, 5xx errors will be thrown.
- Keep concurrency below 15.

Errors and troubleshooting

- 401 Unauthorized: Missing/invalid API key
- 400 Bad Request: Invalid JSON or unsupported parameter type
- 429 Too Many Requests: Rate Limit hit (retry w/ backoff)
- 5xx: transient server errors (retry w/ backoff)